

# lab0-pandas-auto\_mpg

September 22, 2022

## 1 Pandas

As described at <https://pandas.pydata.org> > pandas is an open source, BSD-licensed library providing high-performance, easy-to-use data structures and data analysis tools for the Python programming language.

### 1.1 Resources

1. Ch 5-6 in Python for Data Analysis, 2nd Ed, Wes McKinney (UCalgary library and <https://github.com/wesm/pydata-book>)
2. Ch 3 in Python Data Science Handbook, Jake VanderPlas (Ucalgary library and <https://github.com/jakevdp/PythonDataScienceHandbook>)

Let's explore some of the features.

First, import Pandas, and Numpy as a good companion.

```
[ ]: import numpy as np
import pandas as pd
```

### 1.2 Create pandas DataFrames

There are several ways to create Pandas DataFrames, most notably from reading a csv (comma separated values file). DataFrames are 'spreadsheets' in Python. We will often use `df` as a variable name for a DataFrame.

If data is not stored in a file, a DataFrame can be created from a dictionary of lists

```
data = {'state': ['Ohio', 'Ohio', 'Ohio', 'Nevada', 'Nevada', 'Nevada'],
        'year': [2000, 2001, 2002, 2001, 2002, 2003],
        'pop': [1.5, 1.7, 3.6, 2.4, 2.9, 3.2]}
frame = pd.DataFrame(data)
```

where dictionary keys become column headers.

An alternative is to create from a numpy array and set column headers separately:

```
[ ]: # From a numpy array
df = pd.DataFrame( np.arange(20).reshape(5,4), columns=['alpha', 'beta', 'gamma', 'delta'])
```

```
df
```

```
[ ]:   alpha  beta  gamma  delta
0      0     1      2      3
1      4     5      6      7
2      8     9     10     11
3     12    13     14     15
4     16    17     18     19
```

```
[ ]: # checking its type
type(df)
```

```
[ ]: pandas.core.frame.DataFrame
```

### 1.3 Indexing

Accessing data in Dataframes is done by rows and columns, either index or label based.

```
[ ]: # select a column
df['alpha']
```

```
[ ]: 0      0
1      4
2      8
3     12
4     16
Name: alpha, dtype: int64
```

```
[ ]: # select two columns
df[['alpha', 'gamma']]
```

```
[ ]:   alpha  gamma
0      0      2
1      4      6
2      8     10
3     12     14
4     16     18
```

```
[ ]: # select rows
df.iloc[:2]
```

```
[ ]:   alpha  beta  gamma  delta
0      0     1      2      3
1      4     5      6      7
```

```
[ ]: # select rows and columns
df.iloc[:2, :2]
```

```
[ ]:   alpha  beta
      0      0      1
      1      4      5
```

```
[ ]: # select rows and columns, mixed
      df.loc[:2, ['alpha', 'beta']]
```

```
[ ]:   alpha  beta
      0      0      1
      1      4      5
      2      8      9
```

## 1.4 DataFrame math

Similar to Numpy, DataFrames support direct math

```
[ ]: # direct math
      df2 = (9/5) * df + 32
      df2
```

```
[ ]:   alpha  beta  gamma  delta
      0  32.0  33.8   35.6   37.4
      1  39.2  41.0   42.8   44.6
      2  46.4  48.2   50.0   51.8
      3  53.6  55.4   57.2   59.0
      4  60.8  62.6   64.4   66.2
```

```
[ ]: # add two dataframes of same shape
      df + df2
```

```
[ ]:   alpha  beta  gamma  delta
      0  32.0  34.8   37.6   40.4
      1  43.2  46.0   48.8   51.6
      2  54.4  57.2   60.0   62.8
      3  65.6  68.4   71.2   74.0
      4  76.8  79.6   82.4   85.2
```

```
[ ]: # map a function to each column
      f = lambda x: x.max() - x.min()

      df.apply(f)
```

```
[ ]: alpha    16
      beta     16
      gamma    16
      delta    16
      dtype: int64
```

## 1.5 DataFrame manipulation

Adding and deleting columns, as well as changing entries is similar to Python dictionaries.

Note that most DataFrame methods do not change the DataFrame directly, but return a new DataFrame. It is always good to check how the method you are invoking behaves.

```
[ ]: # add a column
df['epsilon'] = ['low', 'medium', 'low', 'high', 'high']
df
```

```
[ ]:
alpha  beta  gamma  delta  epsilon
0      0     1      2      3      low
1      4     5      6      7      medium
2      8     9     10     11      low
3     12    13     14     15      high
4     16    17     18     19      high
```

```
[ ]: # What is the size?
df.shape
```

```
[ ]: (5, 5)
```

```
[ ]: # delete column
df_dropped = df.drop(columns=['gamma'])
df_dropped
```

```
[ ]:
alpha  beta  delta  epsilon
0      0     1      3      low
1      4     5      7      medium
2      8     9     11      low
3     12    13     15      high
4     16    17     19      high
```

```
[ ]: # the original dataframe is unaffected
df
```

```
[ ]:
alpha  beta  gamma  delta  epsilon
0      0     1      2      3      low
1      4     5      6      7      medium
2      8     9     10     11      low
3     12    13     14     15      high
4     16    17     18     19      high
```

Let's create a copy and assign new values to the first column:

```
[ ]: df_copy = df.copy()
df_copy['alpha'] = 20
print(df)
```

```
print(df_copy)
```

	alpha	beta	gamma	delta	epsilon
0	0	1	2	3	low
1	4	5	6	7	medium
2	8	9	10	11	low
3	12	13	14	15	high
4	16	17	18	19	high

  

	alpha	beta	gamma	delta	epsilon
0	20	1	2	3	low
1	20	5	6	7	medium
2	20	9	10	11	low
3	20	13	14	15	high
4	20	17	18	19	high

DataFrames can be sorted by column:

```
[ ]: # sorting values
df.sort_values(by='epsilon')
```

```
[ ]:
   alpha  beta  gamma  delta  epsilon
3     12    13     14     15     high
4     16    17     18     19     high
0       0     1      2      3      low
2       8     9     10     11      low
1       4     5      6      7     medium
```

## 1.6 Load data from file

Most often data will come from somewhere, often csv files, and using `pd.read_csv()` will allow smooth creation of DataFrames.

Let's load `auto-mpg.csv`:

```
[ ]: data = pd.read_csv('auto-mpg.csv')
```

After loading data, it is good practice to check what we have. Usually, the sequence is: 1. Check dimension 2. Peek at the first rows 3. Get info on data types and missing values 4. Summarize columns

```
[ ]: # Check dimension (rows, columns)
data.shape
```

```
[ ]: (398, 9)
```

```
[ ]: # Peek at the first rows
data.head()
```

```
[ ]:      mpg  cylinders  displacement  horsepower  weight  acceleration  model year  \
0  18.0           8         307.0         130    3504          12.0         70
1  15.0           8         350.0         165    3693          11.5         70
2  18.0           8         318.0         150    3436          11.0         70
3  16.0           8         304.0         150    3433          12.0         70
4  17.0           8         302.0         140    3449          10.5         70

      origin          car name
0         1  chevrolet chevelle malibu
1         1      buick skylark 320
2         1    plymouth satellite
3         1      amc rebel sst
4         1      ford torino
```

```
[ ]: # Column names are
data.columns
```

```
[ ]: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
          'acceleration', 'model year', 'origin', 'car name'],
          dtype='object')
```

```
[ ]: # Get info on data types and missing values
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg             398 non-null   float64
1   cylinders       398 non-null   int64
2   displacement    398 non-null   float64
3   horsepower      398 non-null   object
4   weight          398 non-null   int64
5   acceleration    398 non-null   float64
6   model year      398 non-null   int64
7   origin          398 non-null   int64
8   car name        398 non-null   object
dtypes: float64(3), int64(4), object(2)
memory usage: 28.1+ KB
```

## 1.7 Summarize values

What is the mean, std, min, max in each column?

```
[ ]: data.mean(numeric_only=True)
```

```
[ ]: mpg                23.514573
     cylinders          5.454774
     displacement      193.425879
     weight            2970.424623
     acceleration      15.568090
     model year        76.010050
     origin            1.572864
     dtype: float64
```

```
[ ]: # where are the other columns? Check data types
     data.dtypes
```

```
[ ]: mpg                float64
     cylinders          int64
     displacement      float64
     horsepower        object
     weight            int64
     acceleration      float64
     model year        int64
     origin            int64
     car name          object
     dtype: object
```

Notice that many columns are of type object, which is not a number. Maybe this has to do with missing values? We know from peeking at the first rows that there are '?' values in there. Let's replace these with the string NaN for not-a-number.

```
[ ]: # replace '?' with 'NaN'
     data = data.replace({'?': 'NaN'})
     data.head()
```

```
[ ]:      mpg  cylinders  displacement  horsepower  weight  acceleration  model year  \
0  18.0         8         307.0         130    3504         12.0         70
1  15.0         8         350.0         165    3693         11.5         70
2  18.0         8         318.0         150    3436         11.0         70
3  16.0         8         304.0         150    3433         12.0         70
4  17.0         8         302.0         140    3449         10.5         70

      origin  car name
0         1  chevrolet chevelle malibu
1         1      buick skylark 320
2         1    plymouth satellite
3         1      amc rebel sst
4         1      ford torino
```

Pandas knows that 'NaN' probably means that numbers are missing. Now we can convert the data type from object to float

```
[ ]: # using dictionary to convert specific columns only. We want to keep 'car name'
      ↪as an object
convert_dict = {
    'horsepower': float
}

# convert dtypes
data = data.astype(convert_dict)
data.dtypes
```

```
[ ]: mpg                float64
     cylinders          int64
     displacement       float64
     horsepower         float64
     weight             int64
     acceleration       float64
     model year         int64
     origin             int64
     car name           object
     dtype: object
```

We could have loaded the data with the `na_values` argument to indicate that “?” means missing number:

```
[ ]: data = pd.read_csv('auto-mpg.csv', na_values='?')
     data.dtypes
```

```
[ ]: mpg                float64
     cylinders          int64
     displacement       float64
     horsepower         float64
     weight             int64
     acceleration       float64
     model year         int64
     origin             int64
     car name           object
     dtype: object
```

This worked nicely. Now we can describe all columns, meaning printing basic statistics. Note that by default Pandas ignores NaN, whereas Numpy does not.

```
[ ]: data.describe() # ignores NaN
```

```
[ ]:
count    mpg    cylinders    displacement    horsepower    weight \
mean     23.514573    5.454774    193.425879    104.469388    2970.424623
std       7.815984    1.701004    104.269838    38.491160    846.841774
min       9.000000    3.000000    68.000000    46.000000    1613.000000
```



25%	17.500000	4.000000	104.250000	75.000000	2223.750000
50%	23.000000	4.000000	148.500000	93.500000	2803.500000
75%	29.000000	8.000000	262.000000	126.000000	3608.000000
max	46.600000	8.000000	455.000000	230.000000	5140.000000

	acceleration	model year	origin
count	398.000000	398.000000	398.000000
mean	15.568090	76.010050	1.572864
std	2.757689	3.697627	0.802055
min	8.000000	70.000000	1.000000
25%	13.825000	73.000000	1.000000
50%	15.500000	76.000000	1.000000
75%	17.175000	79.000000	2.000000
max	24.800000	82.000000	3.000000

We could be interested by these statistics in each of the origins. To get these, we first group values by origin, then ask for the description. We will only look at mpg for clarity

```
[ ]: data.groupby(by='origin').describe().mpg
```

```
[ ]:      count      mean      std  min  25%  50%  75%  max
origin
1      249.0  20.083534  6.402892   9.0  15.0  18.5  24.00  39.0
2       70.0  27.891429  6.723930  16.2  24.0  26.5  30.65  44.3
3       79.0  30.450633  6.090048  18.0  25.7  31.6  34.05  46.6
```

## 1.8 Find NaNs

How many NaNs in each column?

We can ask which entries are null, which produces a boolean array

```
[ ]: data.isnull()
```

```
[ ]:      mpg  cylinders  displacement  horsepower  weight  acceleration  \
0    False      False      False      False      False      False
1    False      False      False      False      False      False
2    False      False      False      False      False      False
3    False      False      False      False      False      False
4    False      False      False      False      False      False
..     ...      ...      ...      ...      ...      ...
393  False      False      False      False      False      False
394  False      False      False      False      False      False
395  False      False      False      False      False      False
396  False      False      False      False      False      False
397  False      False      False      False      False      False

      model year  origin  car name
```

```

0      False  False  False
1      False  False  False
2      False  False  False
3      False  False  False
4      False  False  False
..      ...    ...    ...
393    False  False  False
394    False  False  False
395    False  False  False
396    False  False  False
397    False  False  False

```

[398 rows x 9 columns]

Applying `sum()` to this boolean array will count the number of `True` values in each column

```
[ ]: data.isnull().sum()
```

```

[ ]: mpg          0
     cylinders     0
     displacement  0
     horsepower    6
     weight        0
     acceleration  0
     model year    0
     origin        0
     car name      0
     dtype: int64

```

We get complementary information from `info()`

```
[ ]: data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   mpg             398 non-null   float64
 1   cylinders        398 non-null   int64
 2   displacement     398 non-null   float64
 3   horsepower       392 non-null   float64
 4   weight           398 non-null   int64
 5   acceleration     398 non-null   float64
 6   model year       398 non-null   int64
 7   origin           398 non-null   int64
 8   car name         398 non-null   object
dtypes: float64(4), int64(4), object(1)

```

memory usage: 28.1+ KB

We can fill (replace) these missing values, for example with the minimum value in each column

```
[ ]: data.fillna(data.min()).describe()
```

```
[ ]:
      mpg  cylinders  displacement  horsepower  weight \
count  398.000000  398.000000    398.000000    398.000000  398.000000
mean    23.514573    5.454774    193.425879    103.587940  2970.424623
std     7.815984    1.701004    104.269838    38.859575   846.841774
min     9.000000    3.000000     68.000000     46.000000  1613.000000
25%    17.500000    4.000000    104.250000     75.000000  2223.750000
50%    23.000000    4.000000    148.500000     92.000000  2803.500000
75%    29.000000    8.000000    262.000000    125.000000  3608.000000
max    46.600000    8.000000    455.000000    230.000000  5140.000000

      acceleration  model year  origin
count  398.000000  398.000000  398.000000
mean    15.568090    76.010050    1.572864
std     2.757689    3.697627    0.802055
min     8.000000    70.000000    1.000000
25%    13.825000    73.000000    1.000000
50%    15.500000    76.000000    1.000000
75%    17.175000    79.000000    2.000000
max    24.800000    82.000000    3.000000
```

## 1.9 Count unique values (a histogram)

We finish off, with our good friend the histogram

```
[ ]: data['mpg'].value_counts()
```

```
[ ]: 13.0    20
      14.0    19
      18.0    17
      15.0    16
      26.0    14
      ..
      31.9     1
      16.9     1
      18.2     1
      22.3     1
      44.0     1
      Name: mpg, Length: 129, dtype: int64
```